

Launch AWS

Create cluster

ec2-44-191

Clusters | C

MongoDB

practical la

Big Data T

Crimes -

Databricks

+

Ask Gemini

us-east-1.console.aws.amazon.com/emr/home?region=us-east-1#/cloneCluster/j-AIKCQVX20GOD

us-east-1

Search

[Alt+S]

United States (N. Virginia)

Account ID: 3397-1298-6249

voclabs/user5000569=shraddha256@gmail.com

Amazon EMR

EMR on EC2: Clusters

Create cluster

1

sparkcluster

Amazon EMR release [Info](#)
A release contains a set of applications which can be installed on your cluster.
emr-5.20.1

This EMR release reaches End of Support on Aug-01-2026 and will no longer be eligible for technical support. AWS strongly recommends that you run your workloads on the latest Amazon EMR release to receive security-critical updates and fixes. You can also use new Spark upgrade agent to modernize your existing applications on version 5.40 or higher to latest EMR version. To learn more, see [EMR Standard Support policy](#) and [Spark Upgrades](#).

Application bundle

Spark

Core Hadoop

HBase

Presto

Custom

☐ Flink 1.6.2
☐ HCatalog 2.3.4
☒ Hue 4.3.0
☐ MXNet 1.3.1
☐ Phoenix 4.14.0
☒ Spark 2.4.0
☒ Tez 0.9.1

☒ Ganglia 3.7.2
☒ Hadoop 2.8.5
☒ JupyterHub 0.9.4
☐ Mahout 0.13.0
☐ Pig 0.17.0
☐ Sqoop 1.4.7
☒ Zeppelin 0.8.0

☐ HBase 1.4.8
☒ Hive 2.3.4
☐ Livy 0.5.0
☐ Oozie 5.0.0
☐ Presto 0.214
☐ TensorFlow 1.12.0
☐ ZooKeeper 3.4.13

Summary [Info](#)

Name and applications - required

Name
sparkcluster

Amazon EMR release
emr-5.20.1

Application bundle
Custom (Ganglia 3.7.2, Hadoop 2.8.5, Hive 2.3.4, Hue 4.3.0, JupyterHub 0.9.4, Spark 2.4....)

Cluster configuration - required

Uniform instance groups
Primary (m4.large), Core (m4.large)

Cluster scaling and provisioning - required

Provisioning configuration
Core size: 1 instance

Networking - required

Cancel

Clone cluster

CloudShell

Feedback

© 2026, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

35°C
Sunny

Search

ENG
IN

12:11
09-05-2026

Launch AWS

Create cluster

ec2-44-191

Clusters

MongoDB

practical la

Big Data Tr

Crimes

Databricks

Ask Gemini

us-east-1.console.aws.amazon.com/emr/home?region=us-east-1#/cloneCluster/j-AIKCQVX20GOD

aws

Search

[Alt+S]

United States (N. Virginia)

Account ID: 3397-1298-6249

voclabs/user5000569=shraddha256@gmail.com

Amazon EMR

EMR on EC2: Clusters

Create cluster

Uniform instance groups

Choose the same EC2 instance type and purchasing option (On-Demand or Spot) for all nodes in your node group. [Learn more](#)

Provisioning options

Choose from the widest variety of provisioning options for the EC2 instances in your cluster. Diversify instance types and purchasing options, and use an allocation strategy. [Learn more](#)

Uniform instance groups

Primary

Choose EC2 instance type

m4.large

2 vCore 8 GiB memory

EBS only storage On-Demand price: -

Lowest Spot price: -

Actions

☐ Use high availability

Launch highly available, more resilient cluster with three primary nodes on On-Demand Instances. This configuration applies for the lifetime of your cluster. [Learn more](#)

Node configuration - optional

Core

Choose EC2 instance type

m4.large

2 vCore 8 GiB memory

EBS only storage On-Demand price: -

Lowest Spot price: -

Actions

Remove instance group

Node configuration - optional

Summary

Name and applications - required

Name

sparkcluster

Amazon EMR release

emr-5.20.1

Application bundle

Custom (Ganglia 3.7.2, Hadoop 2.8.5, Hive 2.3.4, Hue 4.3.0, JupyterHub 0.9.4, Spark 2.4....)

Cluster configuration - required

Uniform instance groups

Primary (m4.large), Core (m4.large)

Cluster scaling and provisioning - required

Provisioning configuration

Core size: 1 instance

Networking - required

Cancel

Clone cluster

CloudShell

Feedback

© 2026, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

35°C

Sunny

Search

12:11

09-05-2026

Launch AWS x Properties x ec2-44-195 x Clusters | C x MongoDB x practical la x Big Data Tr x Crimes x Databricks x + Ask Gemini

us-east-1.console.aws.amazon.com/emr/home?region=us-east-1#/clusterDetails/j-AIKCQVX20GOD

aws Search [Alt+S] United States (N. Virginia) Account ID: 3397-1298-6249 voclabs/user5000569=shraddha256@gmail.com

Amazon EMR > EMR on EC2: Clusters > sparkcluster

sparkcluster Updated less than a minute ago Terminate Clone in AWS CLI Clone

This EMR release reaches End of Support on Aug-01-2026 and will no longer be eligible for technical support. AWS strongly recommends that you run your workloads on the latest Amazon EMR release to receive security-critical updates and fixes. You can also use new Spark upgrade agent to modernize your existing applications on version 5.40 or higher to latest EMR version. To learn more, see [EMR Standard Support policy](#) and [Spark Upgrades](#).

Summary

Cluster info

Cluster ID
j-AIKCQVX20GOD

Cluster ARN
arn:aws:elasticmapreduce:us-east-1:339712986249:cluster/j-AIKCQVX20GOD

Cluster configuration
Instance groups

Capacity
1 Primary 1 Core 0 Task

Applications

Amazon EMR version
emr-5.20.1

Installed applications
Ganglia 3.7.2, Hadoop 2.8.5, Hive 2.3.4, Hue 4.3.0, JupyterHub 0.9.4, Spark 2.4.0, Tez 0.9.1, Zeppelin 0.8.0

Cluster management

Log destination in Amazon S3
Logging not configured

Primary node public DNS
ec2-44-195-2-208.compute-1.amazonaws.com
[Connect to the Primary node using SSH](#)
[Connect to the Primary node using SSM](#)

Status and time

Status
Waiting

Creation time
May 09, 2026, 11:10 (UTC+05:30)

Elapsed time
1 hour, 1 minute

Properties Bootstrap actions Instances (Hardware) Steps Applications Configurations Monitoring Events Tags (0)

Cluster logs Cluster termination and node replacement

CloudShell Feedback © 2026, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

35°C Sunny Search 12:11 09-05-2026



ec2-44-195-2-208.compute-1.amazonaws.com (hadoop)

Terminal Sessions View X server Tools Games Settings Macros Help

Session Servers Tools Games Sessions View Split MultiExec Tunneling Footages Settings Help

Quick connect...

/home/hadoop/

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: false

```
hive> CREATE DATABASE IF NOT EXISTS chicago_db;
OK
Time taken: 0.653 seconds
hive> USE chicago_db;
OK
Time taken: 0.028 seconds
hive> CREATE EXTERNAL TABLE crimes_raw (
  > crime_id INT,
  > case_number STRING,
  > crime_date STRING,
  > block STRING,
  > primary_type STRING,
  > description STRING,
  > location_description STRING,
  > arrest STRING,
  > domestic STRING,
  > beat INT,
  > district INT,
  > ward INT,
  > community_area INT,
  > crime_year INT,
  > updated_on STRING,
  > latitude DOUBLE,
  > longitude DOUBLE,
  > )
  > ROW FORMAT DELIMITED
  > FIELDS TERMINATED BY ','
  > STORED AS TEXTFILE
  > LOCATION '/user/hadoop/';
NoViableAltException(347@[ ]
    at org.apache.hadoop.hive.ql.parse.HiveParser.columnNameTypeOrPKOrFK(HiveParser.java:33341)
    at org.apache.hadoop.hive.ql.parse.HiveParser.columnNameTypeOrPKOrFKList(HiveParser.java:29513)
    at org.apache.hadoop.hive.ql.parse.HiveParser.createTableStatement(HiveParser.java:6175)
    at org.apache.hadoop.hive.ql.parse.HiveParser.ddlStatement(HiveParser.java:3868)
    at org.apache.hadoop.hive.ql.parse.HiveParser.execStatement(HiveParser.java:2382)
    at org.apache.hadoop.hive.ql.parse.HiveParser.statement(HiveParser.java:1333)
    at org.apache.hadoop.hive.ql.parse.ParseDriver.parse(ParseDriver.java:208)
    at org.apache.hadoop.hive.ql.parse.ParseUtils.parse(ParseUtils.java:77)
```

Remote monitoring

Follow terminal folder


Loading remote monitoring, please wait...

UNREGISTERED VERSION - Please support MobaxTerm by subscribing to the professional edition here: <https://mobaxterm.mobatek.net>

Feels hotter Now

Search

ENG IN 12:49 09-05-2026



ec2-44-195-2-208.compute-1.amazonaws.com (hadoop)

Terminal Sessions View X server Tools Games Settings Macros Help

Session Servers Tools Games Sessions View Split MultiExec Tunneling Footages Settings Help

Quick connect...

/home/hadoop/

Home

aws

ssh

bash_history

bash_profile

bashrc

hivestory

arrest

beat

block

case_number

community_area

crime_date

crime_id

crime_year

description

district

domestic

FIELDS

latitude

LOCATION

location_description

longitude

primary_type

ROW

STORED

updated_on

ward

Remote monitoring

Follow terminal folder

Loading remote monitoring, please wait...

```
000000!0o000Bq-96=:00mE0k$000F#i 000!000!0-0000m0{000d2 0;0N[00];zN!0-#90)$0i00
ULL NULL NULL NULL NULL NULL NULL NULL NULL NULL NULL NULL NULL NULL NULL NULL NULL NULL
Time taken: 2.883 seconds, Fetched: 5 row(s)
hive> SET hive.exec.dynamic.partition=true;
hive> SET hive.exec.dynamic.partition.mode=nonstrict;
hive> CREATE TABLE crimes_partitioned (
> id STRING,
> case_number STRING,
> crime_date STRING,
> block STRING,
> iucr STRING,
> primary_type STRING,
> description STRING,
> location_description STRING,
> arrest BOOLEAN,
> domestic BOOLEAN,
> beat INT,
> district INT,
> ward INT,
> community_area INT,
> fbi_code STRING,
> x_coordinate DOUBLE,
> y_coordinate DOUBLE,
> updated_on STRING,
> latitude DOUBLE,
> longitude DOUBLE,
> location STRING
> )
> PARTITIONED BY (crime_year INT)
> CLUSTERED BY (district) INTO 4 BUCKETS
> ROW FORMAT DELIMITED
> FIELDS TERMINATED BY ','
> STORED AS ORC;
OK
Time taken: 0.134 seconds
hive> SET hive.enforce.bucketing=true;
hive> INSERT INTO TABLE crimes_partitioned
> PARTITION (crime_year)
> SELECT
> id,
```

UNREGISTERED VERSION - Please support Mobaxterm by subscribing to the professional edition here: <https://mobaxterm.mobatek.net>

Feels hotter Now

Search

ENG IN 12:50 09-05-2026



ec2-44-195-2-208.compute-1.amazonaws.com (hadoop)

Terminal Sessions View X server Tools Games Settings Macros Help

Session Servers Tools Games Sessions View Split MultiExec Tunneling Footages Settings Help

Quick connect...

/home/hadoop/

OK

Time taken: 0.2 seconds

```
hive> SET hive.exec.dynamic.partition=true;
hive> SET hive.exec.dynamic.partition.mode=nonstrict;
hive> SET hive.enforce.bucketing=true;
hive> INSERT INTO TABLE crimes_partitioned
  > PARTITION (crime_year)
  > SELECT
    > crime_id,
    > case_number,
    > crime_date,
    > block,
    > primary_type,
    > description,
    > location_description,
    > arrest,
    > domestic,
    > beat,
    > district,
    > ward,
    > community_area,
    > updated_on,
    > latitude,
    > longitude,
    > crime_year
  > FROM crimes_raw;
```

Query ID = hadoop_20260509071618_93a4d1f9-aba4-4ea3-b512-11ff127842e9

Total jobs = 1

Launching Job 1 out of 1

Status: Running (Executing on YARN cluster with App id application_1778305942855_0007)

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	2	2	0	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0	0

VERTICES: 02/02 [=====] 100% ELAPSED TIME: 30.59 s

Loading data to table default.crimes_partitioned partition (crime_year=null)

Remote monitoring

Follow terminal folder

Loading remote monitoring, please wait...

UNREGISTERED VERSION - Please support Mobaxterm by subscribing to the professional edition here: <https://mobaxterm.mobatek.net>

Feels hotter Now

Search

ENG IN 12:50 09-05-2026

ec2-44-195-2-208.compute-1.amazonaws.com (hadoop)

Terminal Sessions View X server Tools Games Settings Macros Help

Session Servers Tools Games Sessions View Split MultiExec Tunneling Footages Settings Help

Quick connect...

/home/hadoop/

▼ Name

- ..
- .aws
- .ssh
- .bash_history
- .bash_profile
- .bashrc
- .jshistory
- arrest
- beat
- block
- case_number
- community_area
- crime_date
- crime_id
- crime_year
- description
- district
- domestic
- FIELDS
- latitude
- LOCATION
- location_description
- longitude
- primary_type
- ROW
- STORED
- updated_on
- ward

Remote monitoring

Follow terminal folder

Loading remote monitoring, please wait...

```
> orrest,
> domestic,
> beat,
> district,
> ward,
> community_area,
> updated_on,
> latitude,
> longitude,
> crime_year
> FROM crimes_raw;
Query ID = hadoop_20260509071618_93a4d1f9-aba4-4ea3-b512-11ff127842e9
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1778305942855_0007)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	2	container	SUCCEEDED	2	2	0	0	0	0
Reducer 2	1	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 02/02 [=====] 100% ELAPSED TIME: 30.59 s
Loading data to table default.crimes_partitioned partition (crime_year=null)

Loaded : 1/1 partitions.
Time taken to load dynamic partitions: 0.425 seconds
Time taken for adding to write entity : 0.003 seconds

OK
Time taken: 36.911 seconds
hive> SELECT COUNT(*) AS total_records
> FROM crimes_partitioned;

OK
473480
Time taken: 0.538 seconds, Fetched: 1 row(s)
hive> SELECT primary_type, COUNT(*) AS total
> FROM crimes_partitioned
> GROUP BY primary_type
> ORDER BY total DESC
```

UNREGISTERED VERSION - Please support Mobaxterm by subscribing to the professional edition here: <https://mobaxterm.mobatek.net>

Feels hotter Now

Search

ENG IN 12:50 09-05-2026

ec2-44-195-2-208.compute-1.amazonaws.com (hadoop)

Terminal Sessions View X server Tools Games Settings Macros Help

Session Servers Tools Games Sessions View Split MultiExec Tunneling Footages Settings Help

Quick connect...

/home/hadoop/

Home

aws

ssh

bash_history

bash_profile

bashrc

hivehistory

arrest

beat

block

case_number

community_area

crime_date

crime_id

crime_year

description

district

domestic

FIELDS

latitude

LOCATION

location_description

longitude

primary_type

ROW

STORED

updated_on

viewd

Remote monitoring

Follow terminal folder

3. ec2-44-195-2-208.compute-1.amazonaws.com

hive> SELECT primary_type, COUNT(*) AS total
> FROM crimes_partitioned
> GROUP BY primary_type
> ORDER BY total DESC
> LIMIT 5;
Query ID = hadoop_20260509071710_d941a0dc-7b8e-4603-82ff-ae1b79e91887
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1778305942855_0007)

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	1	1	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0
Reducer 3	container	SUCCEEDED	1	1	0	0	0	0

VERTICES: 03/03 [=====] 100% ELAPSED TIME: 8.16 s

OK

NULL 229847

"0810" 20114

"0020" 18809

"0486" 18609

"0910" 17160

Time taken: 9.002 seconds, Fetched: 5 row(s)

hive> SELECT crime_year, COUNT(*) AS total_crimes
> FROM crimes_partitioned
> GROUP BY crime_year
> ORDER BY crime_year ASC;
Query ID = hadoop_20260509071725_d422a9c4-6646-4121-ba8a-ad0fcac11d01
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1778305942855_0007)

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	1	1	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0

Loading remote monitoring, please wait...

UNREGISTERED VERSION - Please support MobaxTerm by subscribing to the professional edition here: <https://mobaxterm.mobatek.net>

Feels hotter Now

Search

ENG IN

12:51 09-05-2026

ec2-44-195-2-208.compute-1.amazonaws.com (hadoop)

Terminal Sessions View X server Tools Games Settings Macros Help

Session Servers Tools Games Sessions View Split MultiExec Tunneling Footages Settings Help

Quick connect...

/home/hadoop/

HIVE_DEFAULT_PARTITION 473486
Time taken: 1.976 seconds, Fetched: 1 row(s)
hive> SELECT primary_type, COUNT(*) AS arrest_count
> FROM crimes_partitioned
> WHERE arrest = true
> GROUP BY primary_type
> HAVING COUNT(*) > 500;
Query ID = hadoop_20260509071737_398aa771-2225-468a-b767-f2ace24af782
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1778305942855_0007)

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	1	1	0	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0	0

VERTICES: 02/02 [=====] 100% ELAPSED TIME: 8.70 s

OK
Time taken: 9.545 seconds
hive> SELECT *
> FROM crimes_partitioned
> WHERE crime_year = 2020;
OK
Time taken: 1.222 seconds
hive> EXPLAIN
> SELECT *
> FROM crimes_partitioned
> WHERE crime_year = 2020;
OK
Plan optimized by CBO.
Stage-0
Fetch Operator
Limit:-1
Select Operator [SEL_2]
Output:["_col0","_col1","_col2","_col3","_col4","_col5","_col6","_col7","_col8","_col9","_col10","_col11","_col12","_col13","_col14","_col15","_col16"]
Filter Operator [FIL_4]

Remote monitoring
Follow terminal folder
Loading remote monitoring, please wait...

UNREGISTERED VERSION - Please support Mobaxterm by subscribing to the professional edition here: <https://mobaxterm.mobatek.net>

Feels hotter Now

Search

ENG IN 12:51 09-05-2026

ec2-44-195-2-208.compute-1.amazonaws.com (hadoop)

Terminal Sessions View X server Tools Games Settings Macros Help

Session Servers Tools Games Sessions View Split MultiExec Tunneling Footages Settings Help

Quick connect...

/home/hadoop/

OK

```
> WHERE crime_year = 2020;
Plan optimized by CB0.

Stage-0
Fetch Operator
  limit:-1
  Select Operator [SEL_2]
    Output:["_col0","_col1","_col2","_col3","_col4","_col5","_col6","_col7","_col8","_col9","_col10","_col11","_col12","_col13","_col14","_col15","_col16"]
    Filter Operator [FIL_4]
      predicate:(UDFToDouble(crime_year) = 2020.0)
      TableScan [TS_0]
        Output:["crime_id","case number","crime_date","block","primary_type","description","location_description","arrest","domestic","beat","district","ward",
"d","community_area","updated_on","latitude","longitude"]

Time taken: 0.278 seconds, Fetched: 12 row(s)
hive> SELECT district, COUNT(*) AS domestic_crimes
> FROM crimes_partitioned
> WHERE domestic = true
> GROUP BY district
> ORDER BY domestic_crimes DESC
> LIMIT 3;

Query ID = hadoop_20260509071847_a9e4abb5-df9a-4f5c-83cb-1ecb5007c204
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1778305942855_0007)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	container	SUCCEEDED	1	1	0	0	0	0
Reducer 2	container	SUCCEEDED	1	1	0	0	0	0
Reducer 3	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 03/03 [=====] 100% ELAPSED TIME: 8.09 s
OK
Time taken: 8.829 seconds
hive> [hadoop@ip-172-31-12-237 ~]$ hive
```

Remote monitoring

Follow terminal folder

Loading remote monitoring, please wait...

UNREGISTERED VERSION - Please support Mobaxterm by subscribing to the professional edition here: <https://mobaxterm.mobatek.net>